

Episteme Benchmark Executive Summary

Precision Performance Analytics | Generated: 2026-03-21 02:30:07


Operating System	Linux (amd64)
Java Runtime	25 (Oracle Corporation)
Logical Processors	1

Linear Algebra

Linear Algebra (Analysis)

Linear Algebra | Scaling Performance


Throughput (Operations per Second)



Linear Algebra (Complex HP)


Linear Algebra (Analysis) | Scaling Performance

Throughput (Operations per Second)




Linear Algebra (Dense)

Linear Algebra (Complex HP) | Scaling Performance




Linear Algebra (Distribution Overhead)

Linear Algebra (Dense) | Scaling Performance




Linear Algebra (High-Precision Sparse)

Linear Algebra (Distribution Overhead) | Scaling Performance




Linear Algebra (High-Precision)

Linear Algebra (High-Precision Sparse) | Scaling Performance




Linear Algebra (Inverse)

Linear Algebra (High-Precision) | Scaling Performance




Linear Algebra (Solve)

Linear Algebra (Inverse) | Scaling Performance




Linear Algebra (Sparse)

Linear Algebra (Solve) | Scaling Performance



Linear Algebra (Tensor)

Linear Algebra (Sparse) | Scaling Performance



Linear Algebra (Tensor) | Scaling Performance

Throughput (Operations per Second)

0.000 0.001 0.002 0.003 0.004 0.005 0.006 0.007 0.008 0.009 0.010

[Episteme Native] oneDNN

[Episteme] CUDA Tensor Backend

